

SPS_DATA607_Week3B_DC

David Chen

Window Functions

1. Find (or ask an LLM to generate!) a dataset that includes time series for two or more separate items. For example, you could use end of day stock or cryptocurrency prices since Jan 1, 2022 for several instruments.
2. Use window functions (in SQL or dplyr) to calculate the year-to-date average and the six-day moving averages for each item. ## Approach

I have daily temperature reports for four buildings. Each report contains 24 temperature readings collected every 15 minutes. I plan to combine these CSV files into a single dataset, then calculate the YTD average and a 6-day moving average.

The 6-day moving average will help identify temperature drops, which may indicate boiler issues such as leaks or air vent locks in distributed steam pipes

Data set from <https://www.bitget.com/price/filecoin/historical-data>

Running Code

```
library(dplyr)
```

```
Attaching package: 'dplyr'
```

```
The following objects are masked from 'package:stats':
```

```
filter, lag
```

```
The following objects are masked from 'package:base':
```

```
intersect, setdiff, setequal, union
```

```
library(ggplot2)
library(tidyverse)
```

```
-- Attaching core tidyverse packages ----- tidyverse 2.0.0 --
vforcats    1.0.1      vstringr    1.6.0
v lubridate 1.9.4      v tibble     3.3.1
v purrr      1.2.1      v tidyr     1.3.2
v readr      2.1.6
```

```
-- Conflicts ----- tidyverse_conflicts() --
x dplyr::filter() masks stats::filter()
x dplyr::lag()   masks stats::lag()
i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become non-conflicting
```

```
mydata <- read.csv("https://raw.githubusercontent.com/dyc-sps/SPS_DATA607_Week3B/refs/heads/main.csv")
summary(mydata)
```

timeOpen	timeClose	timeHigh	
Min. :1.535e+12	Min. :1.535e+12	Min. :1.535e+12	
1st Qu.:1.594e+12	1st Qu.:1.594e+12	1st Qu.:1.594e+12	
Median :1.652e+12	Median :1.652e+12	Median :1.652e+12	
Mean :1.653e+12	Mean :1.653e+12	Mean :1.653e+12	
3rd Qu.:1.712e+12	3rd Qu.:1.712e+12	3rd Qu.:1.712e+12	
Max. :1.771e+12	Max. :1.771e+12	Max. :1.771e+12	
timeLow	priceOpen	priceHigh	priceLow
Min. :1.535e+12	Min. : 0.8819	Min. : 0.9422	Min. : 0.6336
1st Qu.:1.594e+12	1st Qu.: 3.5518	1st Qu.: 3.7537	1st Qu.: 3.3225
Median :1.652e+12	Median : 5.0781	Median : 5.3625	Median : 4.7555
Mean :1.653e+12	Mean :15.3725	Mean :16.1979	Mean :14.5478
3rd Qu.:1.712e+12	3rd Qu.:13.0953	3rd Qu.:14.4798	3rd Qu.:11.8922
Max. :1.771e+12	Max. :191.1539	Max. :237.2418	Max. :182.7141
priceClose	volume		
Min. : 0.8819	Min. :5.112e+04		
1st Qu.: 3.5471	1st Qu.:1.433e+07		
Median : 5.0790	Median :1.361e+08		
Mean :15.3672	Mean :2.967e+08		
3rd Qu.:13.0927	3rd Qu.:2.995e+08		
Max. :191.3565	Max. :1.237e+10		

```
head(mydata)
```

```
    timeOpen   timeClose   timeHigh      timeLow priceOpen priceHigh  priceLow
1 1.77072e+12 1.77081e+12 1.77073e+12 1.77079e+12 0.9371375 0.9422195 0.8818362
2 1.77064e+12 1.77072e+12 1.77065e+12 1.77068e+12 0.9327586 0.9459684 0.8937898
3 1.77055e+12 1.77064e+12 1.77056e+12 1.77061e+12 0.9756082 0.9799408 0.9228659
4 1.77047e+12 1.77055e+12 1.77053e+12 1.77049e+12 0.9765257 0.9957679 0.9464217
5 1.77038e+12 1.77047e+12 1.77044e+12 1.77038e+12 0.8819155 0.9852578 0.7998295
6 1.77029e+12 1.77038e+12 1.77030e+12 1.77037e+12 1.0483603 1.0572914 0.8800285
  priceClose   volume
1  0.9004138 81084504
2  0.9371375 92338061
3  0.9327679 81278743
4  0.9756082 137038993
5  0.9765257 254306881
6  0.8818979 247946373
```

```
glimpse(mydata)
```

```
Rows: 2,713
Columns: 9
$ timeOpen   <dbl> 1.77072e+12, 1.77064e+12, 1.77055e+12, 1.77047e+12, 1.77038-
$ timeClose  <dbl> 1.77081e+12, 1.77072e+12, 1.77064e+12, 1.77055e+12, 1.77047-
$ timeHigh   <dbl> 1.77073e+12, 1.77065e+12, 1.77056e+12, 1.77053e+12, 1.77044-
$ timeLow    <dbl> 1.77079e+12, 1.77068e+12, 1.77061e+12, 1.77049e+12, 1.77038-
$ priceOpen  <dbl> 0.9371375, 0.9327586, 0.9756082, 0.9765257, 0.8819155, 1.04-
$ priceHigh  <dbl> 0.9422195, 0.9459684, 0.9799408, 0.9957679, 0.9852578, 1.05-
$ priceLow   <dbl> 0.8818362, 0.8937898, 0.9228659, 0.9464217, 0.7998295, 0.88-
$ priceClose <dbl> 0.9004138, 0.9371375, 0.9327679, 0.9756082, 0.9765257, 0.88-
$ volume     <dbl> 81084504, 92338061, 81278743, 137038993, 254306881, 2479463-
```

Convert the timestamp column to UTC timezone format.

```
mydata <- mydata %>%
  mutate(timeCl = as.Date(as.POSIXct(mydata$timeClose / 1000, origin = "1970-01-01", tz = "UTC"),
  mutate(timeOp = as.Date(as.POSIXct(mydata$timeOpen / 1000, origin = "1970-01-01", tz = "UTC"),
  mutate(timeHi = as.Date(as.POSIXct(mydata$timeHigh / 1000, origin = "1970-01-01", tz = "UTC"),
  mutate(timeLo = as.Date(as.POSIXct(mydata$timeLow / 1000, origin = "1970-01-01", tz = "UTC"),
  colnames(mydata)
```

```

[1] "timeOpen"    "timeClose"   "timeHigh"     "timeLow"      "priceOpen"
[6] "priceHigh"   "priceLow"    "priceClose"   "volume"       "timeCl"
[11] "timeOp"      "timeHi"      "timeLo"

mydata_utc <- mydata %>%
  select(timeopen=timeOp, priceOpen, priceClose, priceHigh, priceLow, volume) %>%
  # select(timeopen=timeOp, timeclose=timeCl, timehigh=timeHi, timelow=timeLo, priceOpen, priceClose)
  mutate(volume_million = round(volume / 1e6, 2))
summary(mydata_utc)

  timeopen            priceOpen            priceClose           priceHigh
Min.   :2018-08-22  Min.   : 0.8819  Min.   : 0.8819  Min.   : 0.9422
1st Qu.:2020-06-30  1st Qu.: 3.5518  1st Qu.: 3.5471  1st Qu.: 3.7537
Median :2022-05-09  Median : 5.0781  Median : 5.0790  Median : 5.3625
Mean   :2022-05-14  Mean   :15.3725  Mean   :15.3672  Mean   :16.1979
3rd Qu.:2024-04-03  3rd Qu.:13.0953  3rd Qu.:13.0927  3rd Qu.:14.4798
Max.   :2026-02-10  Max.   :191.1539  Max.   :191.3565  Max.   :237.2418
               priceLow          volume        volume_million
Min.   : 0.6336  Min.   :5.112e+04  Min.   : 0.05
1st Qu.: 3.3225  1st Qu.:1.433e+07  1st Qu.: 14.33
Median : 4.7555  Median :1.361e+08  Median : 136.07
Mean   :14.5478  Mean   :2.967e+08  Mean   : 296.68
3rd Qu.:11.8922  3rd Qu.:2.995e+08  3rd Qu.: 299.47
Max.   :182.7141  Max.   :1.237e+10  Max.   :12367.98

```

Create year and month fields to support grouping and aggregation.

```

mydata_utc_monthly <- mydata_utc %>%
  mutate(year = year(timeopen), month = month(timeopen, label = TRUE, abbr = TRUE)) %>%
  group_by(year, month)%>%
  summarise(avg_sales = mean(volume_million, na.rm = TRUE), .groups = "drop")

labels <- mydata_utc_monthly %>%
  group_by(year) %>%
  filter(month == max(month)) %>%
  ungroup()

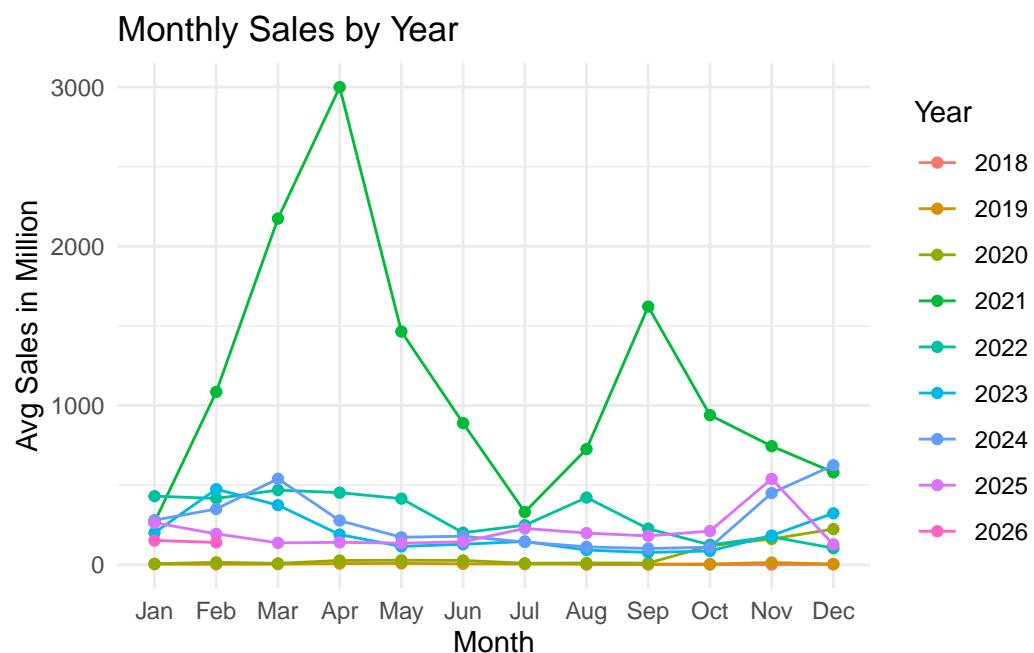
ggplot(mydata_utc_monthly, aes(x = month, y = avg_sales, group = year, color = factor(year)))
  geom_line() +
  geom_point() +
  #geom_text(data = labels, aes(label = year), hjust = -0.2, vjust = 0.5, size = 4) +

```

```

  labs(title = "Monthly Sales by Year",
       x = "Month",
       y = "Avg Sales in Million",
       color = "Year") +
  scale_x_discrete(drop = FALSE) +
  theme_minimal()

```



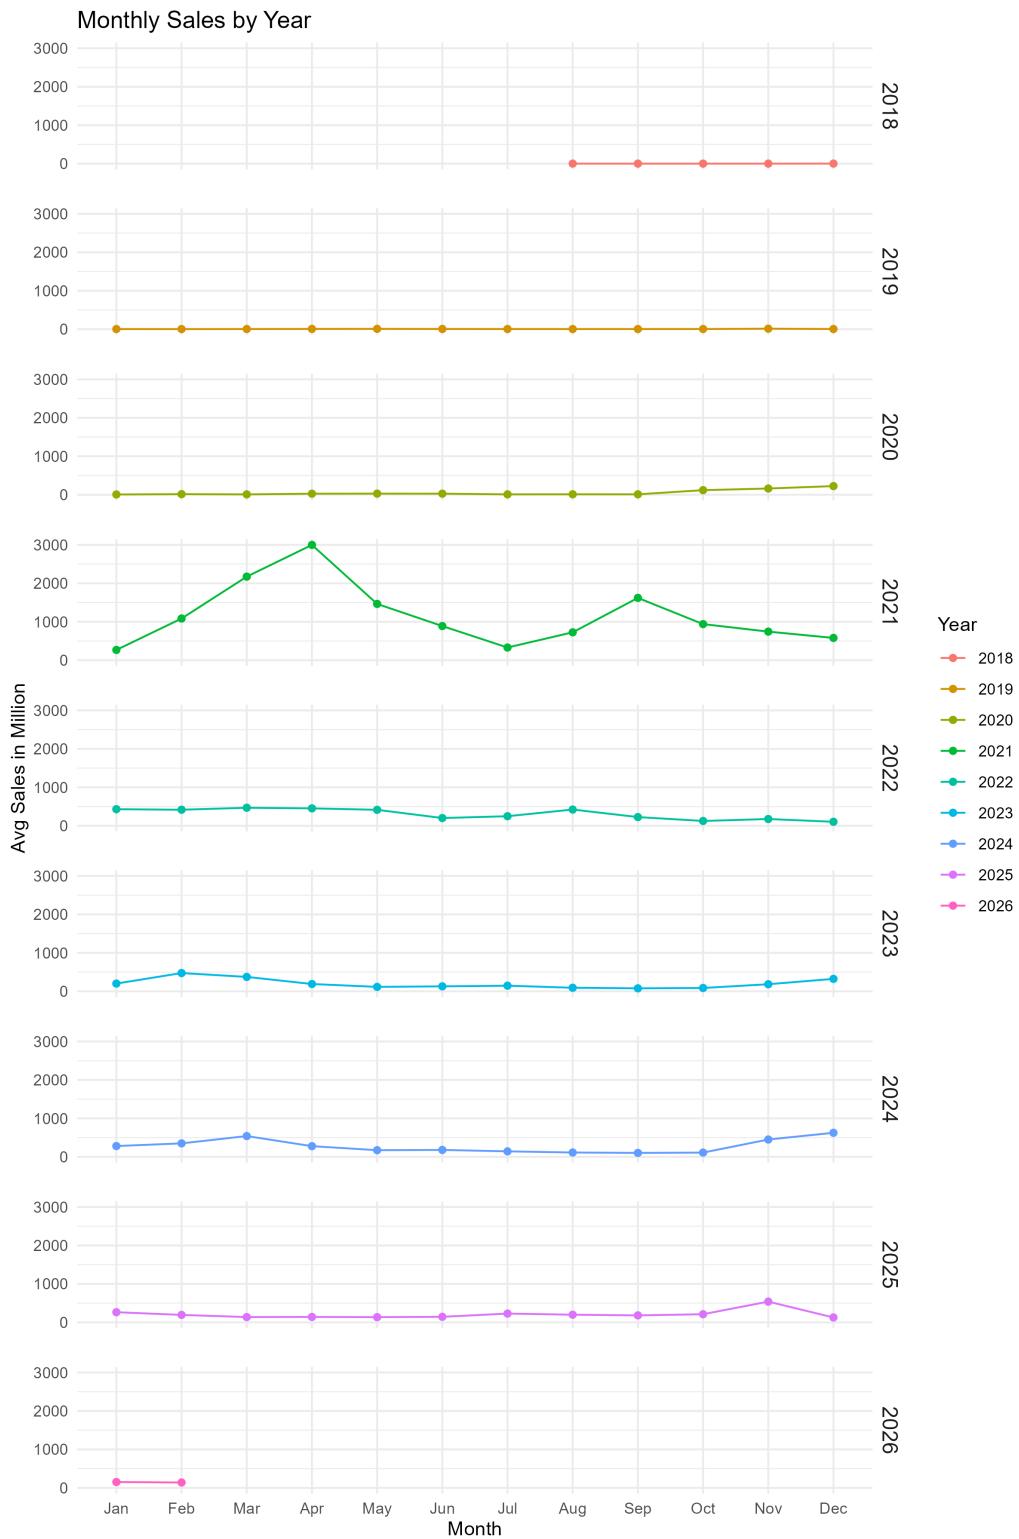
```

ggplot(mydata_utc_monthly, aes(x = month, y = avg_sales, group = year, color = factor(year))) +
  geom_line() +
  geom_point() +
  #geom_text(data = labels,aes(label = year), hjust = -0.2, vjust = 0.5, size = 4) +
  labs(title = "Monthly Sales by Year",
       x = "Month",
       y = "Avg Sales in Million",
       color = "Year") +
  scale_x_discrete(drop = FALSE) +
  theme_minimal() +
  #facet_wrap(~year)
  #facet_wrap(~year, ncol = 1, scales = "free_y")
  #facet_wrap(~year, ncol = 1)

```

```
facet_grid(year ~.)+
  theme(panel.spacing = unit(1.5, "lines"),
        strip.text.y = element_text(size = 12))
```

```
ggsave("monthly_sales.png", width = 8, height = 12)
```



6-days moving average

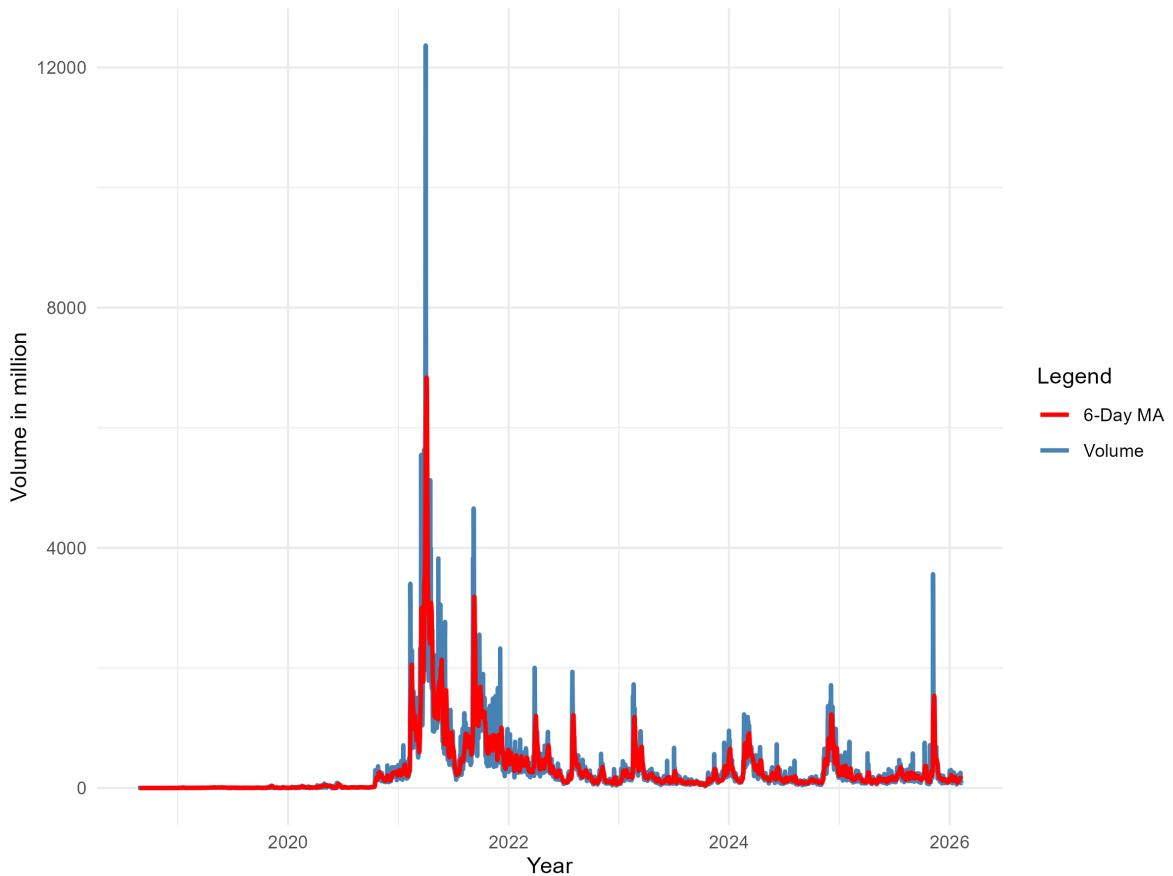
```
#install.packages("slider")
library(slider)
mydata_utc <- mydata_utc %>%
  arrange(timeopen)%>%
  mutate(ma_6day = slide_dbl(volume_million, mean, .before = 5, .complete = TRUE))

ggplot(mydata_utc, aes(x = timeopen)) +
  geom_line(aes(y = volume_million, color = "Volume"), linewidth = 1) +
  geom_line(aes(y = ma_6day, color = "6-Day MA"), linewidth = 1) +
  scale_color_manual(values = c("Volume" = "steelblue", "6-Day MA" = "red")) +
  labs(x = "Year", y = "Volume in million", color = "Legend") +
  theme_minimal()
```

Warning: Removed 5 rows containing missing values or values outside the scale range
(`geom_line()`).

```
ggsave("m6_days_sales.png", width = 8, height = 6)
```

Warning: Removed 5 rows containing missing values or values outside the scale range
(`geom_line()`).



Conclusion

The 6-day moving average, calculated using window functions, provides a smoothed view of the volume over time, reducing the impact of daily fluctuations or anomalies. By applying the window function (slide dbl), each data point incorporates the values of the previous 5 days plus the current day, giving a rolling average that highlights underlying trends.

From the analysis:

1. Short-term spikes or drops in volume are effectively smoothed, making trends more visible. compare the blue line and red line in the last figure
2. Comparing the raw volume to the 6-day moving average allows us to quickly identify periods of sustained increase or decrease.
3. Window functions make it easy to compute such rolling aggregates without collapsing the dataset, preserving the granularity of daily data for further analysis.

Overall, the 6-day moving average serves as a robust indicator of short-term trends, while window functions provide the flexibility to calculate it efficiently and consistently across the dataset.

LLMS used:

- OpenAI. (2025). ChatGPT (Version 5.2) [Large language model]. <https://chat.openai.com>. Accessed Feb 14, 2026.